

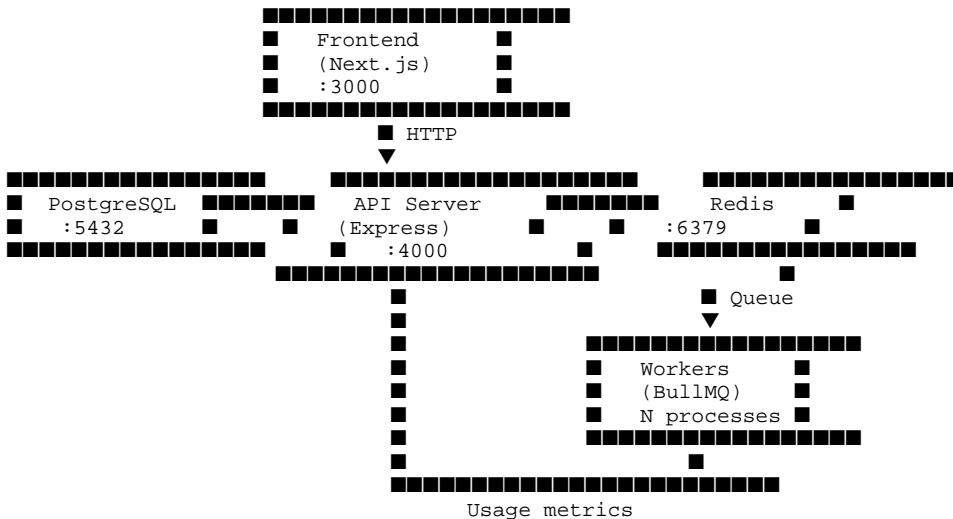
# GPU Cloud Platform

## Architecture Documentation

### 1. System Overview

The platform is a mini AI PaaS that exposes an **OpenAI-compatible** chat completions API behind API key auth, rate limiting, and usage tracking. It uses a **distributed task queue** (Bull/BullMQ + Redis) to distribute inference work across multiple worker processes and stores API keys, deployments, and usage in **PostgreSQL**.

#### Component Diagram



#### Data Flow

**Request** → Client sends POST /v1/chat/completions with Authorization: Bearer <api\_key>.

**API Gateway** → Express: CORS, JSON body, then middleware chain.

**Auth** → authenticateApiKey: load key by prefix, bcrypt compare, attach req.apiKey.

**Rate limit** → Redis sliding window: rate:{apiKeyId}:{window} INCR + EXPIRE; reject if over limit.

**Validation** → validateChatCompletionRequest: model, messages array, roles.

**Inference** → Controller either:

- **Non-streaming**: Call mockInference() in-process, record usage, return JSON.

- **Streaming**: Call mockInference(), then write SSE chunks (data: {...}, data: [DONE]).

**Usage** → recordUsage() inserts into usage\_metrics; recordApiKeyUsage() bumps last\_used\_at and usage\_count on api\_keys.

The **queue** is used when you want to offload work to separate worker processes: jobs are added with addInferenceJob() and workers run Worker('inference', handler). The current HTTP path does in-process mock inference for low latency; the worker is available for scaling out.

## 2. Distributed Design

### Queue pattern

**BullMQ** with Redis: one queue 'inference', job data includes requestId, model, messages, apiKeyId, stream.

**Workers:** 2–3 (or more) Node processes run the same worker script; each pulls jobs from the queue. Concurrency per worker is configurable (e.g. 4).

**Retries:** 3 attempts with exponential backoff (1s, 2s, 4s). Failed jobs can be moved to a dead-letter flow by listening to 'failed' and re-adding elsewhere.

**Idempotency:** Job ID = requestId so duplicate submissions don't create duplicate jobs.

### Worker coordination

No leader election: all workers are equal consumers.

Redis handles atomic pop and locking; BullMQ manages retries and delays.

For 'wait for result' in HTTP, you'd typically use job.getState() / job.waitUntilFinished() or a separate result store (e.g. Redis by requestId) that the worker writes to and the API polls or subscribes to.

### Failure handling

**DB down:** Health check fails; API returns 503 on /health; requests that need DB (auth, usage) fail with 500.

**Redis down:** Rate limiter fails open (next()); queue health fails; job adds would throw.

**Worker crash:** BullMQ marks job failed and retries per policy; after max attempts, job stays in failed set for inspection or DLQ.

## 3. Scalability

### Current limits (single box)

**~100 concurrent requests** if each holds a DB connection and does in-process inference.

**API key validation** <100 ms with indexed key\_prefix and bcrypt compare.

**Queue latency** <500 ms for Redis and a few workers on the same host.

### Scaling to 1000 workers / high RPS

**Redis:** Move to **Redis Cluster** or a managed Redis with high connection and throughput limits; keep queue and rate-limit keys sharded or single-queue per region.

**PostgreSQL:** **Read replicas** for key validation and metrics reads; primary for writes (keys, usage). Connection pooling (e.g. PgBouncer) in front.

**Workers:** Run 1000 worker processes across many nodes; all consuming the same queue(s). Optionally **multiple queues** (e.g. by region or model) to isolate blast radius and prioritize.

**API:** Horizontal scaling of Express behind a **load balancer**; stateless, so no sticky sessions. Health checks for LB to remove bad nodes.

**Frontend:** CDN for static assets; optional serverless or cached dashboard that reads from the same API.

## Geographic distribution

Deploy API + workers + Redis + DB per region. Use **global load balancing** (e.g. latency-based) to route to nearest region.

Cross-region: replicate DB (e.g. read replicas in other regions) or keep usage regional and aggregate in a central analytics DB.

## 4. Trade-offs

Decision	Choice	Why
Queue	Bull/BullMQ	Simple, Redis-based, good for 10k–100k jobs/day; no need for Kafka's throughput and ordering guarantees.
DB	PostgreSQL	ACID, good for keys and time-series usage; simple to run and replicate. DynamoDB would be more complex.
Auth	API key only	Fits "server-to-server" and SDK usage; no OAuth/SSO in scope.
Inference	Mock in-process	Simulates GPU path; swap mockInference() for real model calls (local or remote).
Rate limit	Redis counter	Sliding window per key; sufficient for per-key limits. For global limits, add a separate global counter.

## 5. Future Work

**Multi-tenancy:** Namespace keys and usage by tenant\_id; enforce quotas per tenant.

**Cost tracking:** Store cost per token or per model; aggregate in usage and expose in dashboard and billing.

**Auto-scaling:** Scale worker count (e.g. K8s HPA) based on queue depth or latency.

**Real GPU:** Replace mockInference() with calls to GPU inference services (e.g. Triton, vLLM, or external APIs) and optional streaming from worker to API via Redis or WebSocket.